

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 from matplotlib import pyplot as plt
        4 import seaborn as sns
        5
```

```
In [2]: 1 import warnings
        2 warnings.filterwarnings('ignore')
```

```
In [3]: 1 salary_data=pd.read_csv('Salary_Data.csv')
        2 salary_data
```

Out[3]:

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0
24	8.7	109431.0

	YearsExperience	Salary
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

In [4]: 1 salary_data.isnull().sum()

Out[4]: YearsExperience 0
Salary 0
dtype: int64

In [5]: 1 salary_data.dtypes

Out[5]: YearsExperience float64
Salary float64
dtype: object

In [6]: 1 salary_data.shape

Out[6]: (30, 2)

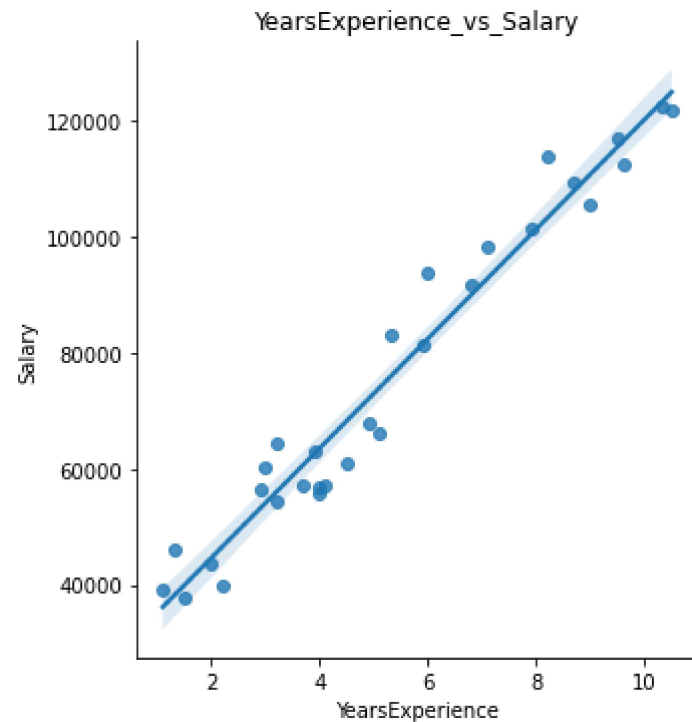
In [7]: 1 salary_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes
```

Assumption check


1. linearity check

```
In [8]: 1 sns.lmplot(x='YearsExperience',y='Salary',data=salary_data)
        2 plt.title('YearsExperience_vs_Salary')
        3 plt.show()
```




2. distribution check

```
In [9]: 1 sns.distplot(a=salary_data['YearsExperience'],hist=True)  
2 plt.show()
```



```
In [10]: 1 sns.distplot(a=salary_data['Salary'],hist=True)  
2 plt.show()
```




```
In [11]: 1 salary_data.corr()
```

Out[11]:

	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

```
In [12]: 1 sns.heatmap(salary_data.corr(),annot=True)  
2 plt.show()
```



model building.|| model training

```
In [13]: 1 import statsmodels.formula.api as smf
```

```
In [14]: 1 linear_model=sfm.ols(formula='Salary~YearsExperience',data=salary_data).fit()
        2 linear_model
```

```
Out[14]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x1a0255e0f70>
```

```
In [15]: 1 linear_model.params
```

```
Out[15]: Intercept          25792.200199
        YearsExperience      9449.962321
        dtype: float64
```

```
In [16]: 1 linear_model.pvalues
```

```
Out[16]: Intercept          5.511950e-12
        YearsExperience      1.143068e-20
        dtype: float64
```

model testing

```
In [17]: 1 y=(3*9449.962321)+ 25792.200199 #manual prediction for 3 years of experience
        2 y
```

```
Out[17]: 54142.087162
```

```
In [18]: 1 z=(4*9449.962321)+ 25792.200199 #manual prediction for 4 years of experience
        2 z
```

```
Out[18]: 63592.049483
```

```
In [19]: 1 data=pd.Series([3,4])# auto prediction for 3 and 4 years of experience
        2 data
```

```
Out[19]: 0    3
        1    4
        dtype: int64
```



```
In [20]: 1 data_pred=pd.DataFrame(data,columns=['YearsExperience'])
         2 data_pred
```

```
Out[20]:
```

	YearsExperience
0	3
1	4

```
In [21]: 1 model_pred=linear_model.predict(data_pred)
         2 model_pred
```

```
Out[21]: 0    54142.087163
         1    63592.049484
         dtype: float64
```

model evaluation

```
In [22]: 1 linear_model=smf.ols(formula='Salary~YearsExperience',data=salary_data).fit()
         2 print('R-square                : ',round(linear_model.rsquared,4))
         3 print('Adjusted R-square       : ',round(linear_model.rsquared_adj,4))
         4 print('Akaike information criterion (AIC) : ',round(linear_model.aic,4))
         5 print('Bayesian information criterion(BIC): ',round(linear_model.bic,4))
```

```
R-square                : 0.957
Adjusted R-square       : 0.9554
Akaike information criterion (AIC) : 606.8823
Bayesian information criterion(BIC): 609.6847
```

```
In [37]: 1 plt.figure(figsize = (8,6))
          2 sns.scatter(x = linear_model["YearsExperience"], y = linear_model["Salary"],data=salary_data)
          3 plt.title("Hetroscedasticity", fontweight = 'bold', fontsize = 14)
          4 plt.show()
```

```
-----
AttributeError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_752\4037791447.py in <module>
      1 plt.figure(figsize = (8,6))
----> 2 sns.scatter(x = linear_model["YearsExperience"], y = linear_model["Salary"],data=salary_data)
      3 plt.title("Hetroscedasticity", fontweight = 'bold', fontsize = 14)
      4 plt.show()
```

AttributeError: module 'seaborn' has no attribute 'scatter'

<Figure size 576x432 with 0 Axes>

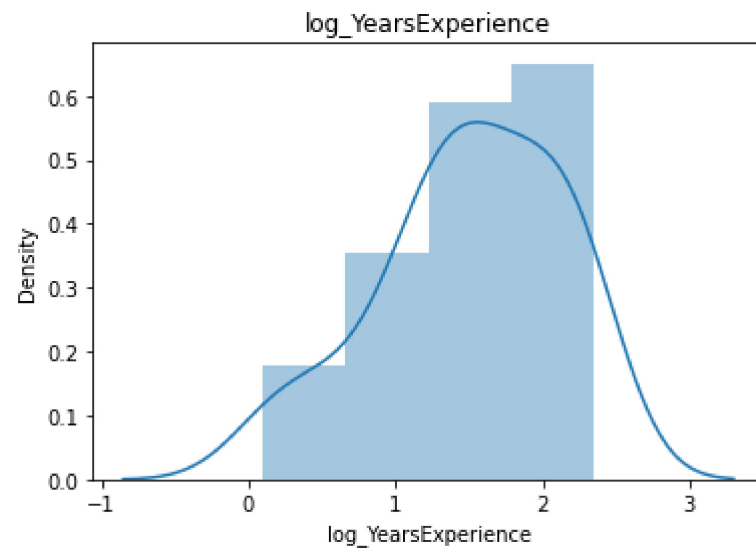
1.Log Transformation

```
In [23]: 1 salary_data['log_YearsExperience']=np.log(salary_data['YearsExperience'])  
2 salary_data.head(10)
```

```
Out[23]:
```

	YearsExperience	Salary	log_YearsExperience
0	1.1	39343.0	0.095310
1	1.3	46205.0	0.262364
2	1.5	37731.0	0.405465
3	2.0	43525.0	0.693147
4	2.2	39891.0	0.788457
5	2.9	56642.0	1.064711
6	3.0	60150.0	1.098612
7	3.2	54445.0	1.163151
8	3.2	64445.0	1.163151
9	3.7	57189.0	1.308333

```
In [24]: 1 sns.distplot(a=salary_data['log_YearsExperience'])  
2 plt.title('log_YearsExperience')  
3 plt.show()
```



```
In [25]: 1 model_1=smf.ols(formula='Salary~log_YearsExperience', data=salary_data).fit()
        2 model_1.summary()
```

Out[25]: OLS Regression Results

Dep. Variable:	Salary	R-squared:	0.854
Model:	OLS	Adj. R-squared:	0.849
Method:	Least Squares	F-statistic:	163.6
Date:	Fri, 01 Jul 2022	Prob (F-statistic):	3.25e-13
Time:	12:09:50	Log-Likelihood:	-319.77
No. Observations:	30	AIC:	643.5
Df Residuals:	28	BIC:	646.3
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.493e+04	5156.226	2.895	0.007	4365.921	2.55e+04
log_YearsExperience	4.058e+04	3172.453	12.792	0.000	3.41e+04	4.71e+04

Omnibus:	1.094	Durbin-Watson:	0.512
Prob(Omnibus):	0.579	Jarque-Bera (JB):	0.908
Skew:	0.156	Prob(JB):	0.635
Kurtosis:	2.207	Cond. No.	5.76

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [26]: 1 model_1.rsquared
```

Out[26]: 0.8538888828756969

This r-square value is less than the r-square of the model from raw data.


2. SQRT MODEL

```
In [27]: 1 salary_data['sqrt_YearsExperience']=np.sqrt(salary_data['YearsExperience'])  
        2 salary_data.head(10)
```

```
Out[27]:
```

	YearsExperience	Salary	log_YearsExperience	sqrt_YearsExperience
0	1.1	39343.0	0.095310	1.048809
1	1.3	46205.0	0.262364	1.140175
2	1.5	37731.0	0.405465	1.224745
3	2.0	43525.0	0.693147	1.414214
4	2.2	39891.0	0.788457	1.483240
5	2.9	56642.0	1.064711	1.702939
6	3.0	60150.0	1.098612	1.732051
7	3.2	54445.0	1.163151	1.788854
8	3.2	64445.0	1.163151	1.788854
9	3.7	57189.0	1.308333	1.923538

```
In [28]: 1 sns.distplot(a=salary_data['sqrt_YearsExperience'])  
2 plt.title('sqrt_YearsExperience')  
3 plt.show()
```



```
In [29]: 1 model_2=smf.ols(formula='Salary~sqrt_YearsExperience', data=salary_data).fit()
        2 model_2.summary()
```

Out[29]: OLS Regression Results

Dep. Variable:	Salary	R-squared:	0.931
Model:	OLS	Adj. R-squared:	0.929
Method:	Least Squares	F-statistic:	377.8
Date:	Fri, 01 Jul 2022	Prob (F-statistic):	8.57e-18
Time:	12:09:50	Log-Likelihood:	-308.52
No. Observations:	30	AIC:	621.0
Df Residuals:	28	BIC:	623.8
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-1.606e+04	4921.599	-3.262	0.003	-2.61e+04	-5974.331
sqrt_YearsExperience	4.15e+04	2135.122	19.437	0.000	3.71e+04	4.59e+04

Omnibus:	0.588	Durbin-Watson:	1.031
Prob(Omnibus):	0.745	Jarque-Bera (JB):	0.638
Skew:	0.011	Prob(JB):	0.727
Kurtosis:	2.286	Cond. No.	9.97

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [30]: 1 model_2.rsquared
```

Out[30]: 0.9310009544993526

This r-square value is less than the r-square of the model from raw data.

3. RECIPROCAL

```
In [31]: 1 salary_data['res_proc_Salary']=1/salary_data['Salary']
          2 salary_data.head(10)
```

```
Out[31]:
```

	YearsExperience	Salary	log_YearsExperience	sqrt_YearsExperience	res_proc_Salary
0	1.1	39343.0	0.095310	1.048809	0.000025
1	1.3	46205.0	0.262364	1.140175	0.000022
2	1.5	37731.0	0.405465	1.224745	0.000027
3	2.0	43525.0	0.693147	1.414214	0.000023
4	2.2	39891.0	0.788457	1.483240	0.000025
5	2.9	56642.0	1.064711	1.702939	0.000018
6	3.0	60150.0	1.098612	1.732051	0.000017
7	3.2	54445.0	1.163151	1.788854	0.000018
8	3.2	64445.0	1.163151	1.788854	0.000016
9	3.7	57189.0	1.308333	1.923538	0.000017


```
In [32]: 1 sns.distplot(a=salary_data['res_proc_Salary'])  
2 plt.title('res_proc_Salary')  
3 plt.show()
```



```
In [33]: 1 model_3=smf.ols(formula='res_proc_Salary~YearsExperience', data=salary_data).fit()
          2 model_3.summary()
```

Out[33]: OLS Regression Results

Dep. Variable:	res_proc_Salary	R-squared:	0.861
Model:	OLS	Adj. R-squared:	0.856
Method:	Least Squares	F-statistic:	173.2
Date:	Fri, 01 Jul 2022	Prob (F-statistic):	1.63e-13
Time:	12:09:50	Log-Likelihood:	350.83
No. Observations:	30	AIC:	-697.7
Df Residuals:	28	BIC:	-694.9
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.454e-05	8.2e-07	29.913	0.000	2.29e-05	2.62e-05
YearsExperience	-1.799e-06	1.37e-07	-13.162	0.000	-2.08e-06	-1.52e-06

Omnibus:	1.760	Durbin-Watson:	1.137
Prob(Omnibus):	0.415	Jarque-Bera (JB):	1.380
Skew:	0.516	Prob(JB):	0.502
Kurtosis:	2.802	Cond. No.	13.2

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [34]: 1 model_3.rsquared
```

Out[34]: 0.8608672473082564

This r-square value is less than the r-square of the model from raw data.

Model Selection

Now by comparing r-square of all models,

we can say that the models which are fitted by using transformation are not so good as compare to our model from raw data(original data)

Hence , we select our first model for further calculation